

Theoretical Report: Natural Language Understanding in Dialogue Systems

Ciobanu Sergiu-Tudor, Moga Patricia

July 31, 2026

Abstract

Dialogue systems are platforms that allow people and computers to interact using natural language. A key part of these systems is Natural Language Understanding (NLU), which works by taking unstructured human input and organizing it into a format the machine can understand. NLU is important because it identifies the user's intent and extracts the specific details needed to handle the request properly, which makes the conversation feel more reliable. This report explores how NLU has evolved by looking at four different models: *Intent Classification and Slot Tagging*, *Sequence-to-sequence architectures*, *JointBERT*, and *TOD-BERT*. By comparing these methods, we show how technology is moving from traditional sequential pipelines toward integrated end-to-end models that are much more flexible for real-world interactions.

Contents

1	Introduction to Dialogue Systems (Ciobanu Sergiu-Tudor)	3
1.1	Conceptual Framework	3
1.2	Application Domains	3
1.3	Basic Dialogue Systems Types	3
1.4	Communication Domains	4
1.5	Modes of Interaction	4
1.6	Foundation and Historical Context	4
2	Natural Language Understanding (Ciobanu Sergiu-Tudor & Moga Patricia)	5
2.1	Challenges in Spoken Language Understanding	5
2.2	Semantic Representations and Methodologies	6
2.3	Shallow Parsing: Classification and Sequence Tagging	6
3	Comparison of NLU Methods (Ciobanu Sergiu-Tudor & Moga Patricia)	7
3.1	Intent Classification and Slot Tagging (Moga Patricia)	7
3.2	Seq2seq-based NLU (Moga Patricia)	9
3.3	Overview of BERT (Ciobanu Sergiu-Tudor)	10
3.4	JointBERT: Joint Intent Classification and Slot Filling (Ciobanu Sergiu-Tudor)	11
3.5	TOD-BERT: Pre-trained NLU for Task-Oriented Dialogue (Ciobanu Sergiu-Tudor)	12
4	Evaluating NLU Performance (Moga Patricia)	13
4.1	Quantitative Metrics	13
4.2	Comparative Benchmarks and Simulation	14
4.3	Insights and Robustness	14
5	Conclusion and Future Directions	15

1 Introduction to Dialogue Systems (Ciobanu Sergiu-Tudor)

Dialogue systems are AI-driven computational platforms designed to facilitate interaction between humans and machines through natural language. By utilizing conversational agents and dialogue management techniques, these systems aim to provide intuitive communication and efficient user experiences.

1.1 Conceptual Framework

A dialogue system is defined as a computer system designed to interact with users in natural language, frequently in spoken form. These systems, often referred to as conversational agents or chatbots, rely on integrated Natural Language Processing (NLP) methods to function.

A Spoken Dialogue Model (SDM) specifically processes and generates multi-turn, human-like spoken interactions. These models are implemented through either cascaded pipelines or end-to-end architectures. They leverage various technical domains, including semantic parsing, Large Language Models (LLMs), reinforcement learning, multi-modal learning, and comprehensive evaluation methodologies.

1.2 Application Domains

The scope of dialogue systems is extensive, encompassing a diverse list of functional environments and specialized technologies. In the consumer sector, they are most prevalent as integrated assistants within mobile operating systems and in-car interfaces, where they provide hands-free control and information retrieval. Beyond personal devices, these systems are frequently deployed in service-oriented roles, such as automated phone hotlines and customer support infrastructures, to streamline high-volume interactions. Furthermore, the technology extends into assistive and social domains, providing essential support in therapy, elderly care, and companionship. This versatility also encompasses the entertainment industry, where dialogue systems power non-player characters (NPCs) in video games and social chatbots.

1.3 Basic Dialogue Systems Types

Dialogue systems are primarily categorized by their functional objectives, specifically distinguishing between task-oriented and non-task-oriented architectures. Task-oriented systems are engineered to complete specific user goals, such as booking transportation, managing restaurant reservations, or controlling smart home devices. Historically, these systems represent the majority of practical applications and are characterized by their reliance on "backend access" to external databases. In contrast, non-task-oriented sys-

tems, or "chitchat" agents, focus on social conversation and entertainment. Their primary success metric is user engagement, often achieved through the simulation of a specific persona or by attempting to satisfy the criteria of the Turing test.

1.4 Communication Domains

The scope of a system is further defined by its conversation domain, which refers to the specific area of interest or topic coverage. Traditional systems are typically closed-domain, restricted to a single, well-defined area such as a specialized banking service. Modern implementations, such as Siri or Alexa, often utilize multi-domain architectures by aggregating several single-domain systems into a unified interface. Beyond these lies the open-domain category, which aims to provide responses across any subject. While originally limited to social chitchat, the advent of Large Language Models has significantly expanded the feasibility of open-domain task handling.

1.5 Modes of Interaction

The complexity of a dialogue system is also dictated by its mode of communication, ranging from text-based interfaces to sophisticated multimodal systems. Text remains the most robust and established medium due to its ease of implementation, though it lacks the full nuance of natural interaction. Voice-based systems offer a more natural experience by conveying emotion, tone, and personality while enabling hands-free utility over telephonic or ambient interfaces. Multimodal systems represent the highest level of complexity, integrating voice and text with visual data such as graphics, gestures, and facial mimics. It is important to note that spoken language is fundamentally distinct from written text. Natural speech is often characterized by ungrammatical structures, restarts, hesitations, and real-time corrections. Furthermore, spoken dialogue involves acoustic variables such as pitch, stress, accents, and dialects, as well as conversational overlaps, all of which increase the difficulty of accurate processing compared to standardized written input.

1.6 Foundation and Historical Context

This report specifically addresses the early stages of Spoken Dialogue Models, with a focus on Statistical Dialogue Systems. Historically, these systems are rooted in modular, cascaded architectures that function as a sequential pipeline. This framework couples several specialized components: Automatic Speech Recognition (ASR), Natural Language Understanding (NLU), Dialogue Management (DM), Natural Language Generation (NLG), and Text-to-Speech (TTS) [6]. In these traditional models, each speech act is treated as a distinct processing stage, with information flowing across rigid boundaries.




Figure 1: The dialog system architecture

The pipeline, as shown in the Figure 1 above, initiates with the ASR module, which captures the user’s acoustic signal and transforms it into a textual hypothesis. This text is then processed by the NLU component, which parses the hypothesis to generate a structured semantic representation of the utterance. This representation serves as the input for the Dialogue Manager, the central component responsible for deciding the system’s next action. The DM maintains the dialogue history and tracks the user profile to ensure contextual relevance. Once an action is determined, the NLG module, functioning essentially as an inverse of the NLU, converts the system’s semantic intent back into a natural language text string, often including prosodic markup. Finally, the TTS engine synthesizes this text into a speech signal to be delivered to the user.

2 Natural Language Understanding (Ciobanu Sergiu-Tudor & Moga Patricia)

Natural Language Understanding (NLU), frequently referred to as spoken language understanding or semantic encoding, serves as the critical bridge between raw human input and machine logic. Its primary objective is to map a sequence of words into a structured, explicit representation of meaning. In the context of a dialogue pipeline, NLU transforms the textual hypothesis provided by the ASR into a format that the Dialogue Manager can process to execute specific tasks.

2.1 Challenges in Spoken Language Understanding

The task of mapping words to meaning is complicated by the inherent noise in spoken communication. Unlike written text, spoken input is often non-grammatical. For in-

stance, a user might provide a fragmented request such as "find something cheap for kids," which the system must still resolve. Furthermore, disfluencies are common, including hesitations (fillers like "uhm" or "uh"), repetitions, and fragments where a thought is left incomplete.

Users also frequently perform self-repairs, such as "find something cheap, no, I mean moderate," requiring the NLU to distinguish between the error and the correction. These linguistic challenges are further compounded by external factors, such as ASR transcription errors, the use of synonyms, and out-of-domain utterances that fall outside the system's predefined capabilities.

2.2 Semantic Representations and Methodologies

NLU systems utilize various frameworks to represent meaning, ranging from hierarchical trees to flat structures. Syntax and semantic trees are common in standalone parsing, while frames provide a way to organize information into slots, though they are not always directly mapped to specific words.

In contemporary dialogue systems, the most prevalent representation is the **Dialogue Act (DA)**, which consists of an intent combined with a set of slots and values. This is typically a flat representation without a complex hierarchy, making it highly efficient for task-oriented systems.

2.3 Shallow Parsing: Classification and Sequence Tagging

Modern NLU focuses predominantly on shallow parsing, which simplifies the task into classification or sequence tagging or a combination of both.

NLU as Classification: In this approach, dialogue acts are treated as a set of semantic concepts. Each concept (intent or slot-value pair) is handled through binary classification: determining if concept Y is present in utterance X , one classifier being used per concept. However, this can lead to consistency issues, such as the system simultaneously detecting conflicting intents (e.g., "affirm" and "negate"). These conflicts must be resolved externally, often by comparing classifier confidence scores.

Slot Filling as Sequence Tagging: To extract specific values without manual delexicalization (replacement of token with placeholders indicating entity type), NLU often employs sequence tagging, similar to Named Entity Recognition (NER). Each word is classified using the IOB (Inside-Outside-Beginning) format. Therefore, Beginning (B) indicates the start of the slot value, Inside (I) denotes the continuation of the slot value, and Outside (O) suggests words that carry no specific slot information.

3 Comparison of NLU Methods (Ciobanu Sergiu-Tudor & Moga Patricia)

The evolution of NLU has resulted in a diverse range of architectural frameworks, including Convolutional Neural Networks (CNNs), Long Short-Term Memory (LSTM) networks, and encoder-decoder Seq2Seq structures enhanced by attention mechanisms. These models are designed to handle core tasks such as intent classification, sequence tagging, and sequence prediction. To improve performance, contemporary approaches often utilize joint models that simultaneously address intent detection and slot filling within a single framework. Furthermore, Transformer-based architectures like BERT have become specialized for NLU, providing advancements in task-oriented dialogue through specialized pre-training techniques.

3.1 Intent Classification and Slot Tagging (Moga Patricia)

The classical approach to Natural Language Understanding operates on the fundamental principle of mapping unstructured words into a structured semantic representation, a process often described as semantic decoding. This method views the NLU task as a way of decomposing meaning into two distinct components: the intent and the slot-value pairs. In this framework, the primary goal is to transform a raw user input into a formal dialogue act that the rest of the system can process. The intent represents the broader objective of the user, while the slots act as specific variables or details that fill in the necessary information for that objective. As dialogue systems have advanced in their ability to understand more of what a user says, it has become increasingly clear that the sophistication of the system's response must keep pace with the quality of its understanding to maintain a natural flow [9].

The technical execution of this pipeline is built on a dual-task strategy that concurrently analyzes two distinct layers of meaning. First, the system identifies the user's intent by evaluating the entire input to determine the overall objective. This is approached as a classification process where the model weighs various predefined categories to select the most likely goal. For example, a request like "Find me a flight to London" would be classified under a "book flight" intent. Simultaneously, the system performs slot tagging to identify and extract specific entities, such as dates, locations, or names, which are essential for completing the semantic frame [5]. This parallel processing of intent and slots is visually represented in Figure 2, which illustrates how the input is split into separate paths to form the final structured output.




Figure 2: The modular architecture of a traditional NLU system showing the parallel processing of intent and slots.

Description of Figure 2: This figure provides a visual representation of how semantic decoding works within a classical NLU pipeline. The process begins with a raw user input, in this case, a request to find a flight to Paris. The diagram illustrates the dual-task strategy by showing how the system splits the analytical effort into two parallel workflows. The top path focuses on the Intent Classifier, which evaluates the entire sentence to produce the label *find_flight*. Meanwhile, the bottom path describes the Slot Tagging process, which fragments the text to isolate key entities like Prague and Paris, labeling them as *origin* and *destination* respectively. Finally, these data streams converge to form a structured Dialogue Act, which allows the dialogue management system to process the information logically.

While this modular design provides a clear structure, its implementation has evolved from rule-based systems to neural classifiers. Modern architectures often treat intent detection as a multi-class classification problem using a softmax function, while slot tagging is handled through sequence tagging or binary classifiers. By using attention, the model can compute a context vector as a weighted average of hidden states, allowing it to focus on specific, relevant parts of the input sequence regardless of its length. This helps the system overcome long-term memory challenges and refer to arbitrary locations in the text to more accurately identify slots. [8]

Despite these neural improvements, the classical pipeline still faces challenges. As new data is introduced, new intents can emerge that semantically overlap with existing ones, potentially weakening the model precision [1]. Additionally, the variety of real-world names often leads to a lack of data. To reduce this problem, delexicalization is used to replace specific entities with generic placeholders, such as *destination*. While this helps the model learn linguistic patterns, traditional pipelines remain relatively rigid compared

to the more flexible, modern neural architectures designed for complex conversational AI [4].

3.2 Seq2seq-based NLU (Moga Patricia)

While the classification-based methods discussed previously are effective, the usual problem is when they encounter words that were not part of their training data, a challenge known as the out-of-vocabulary (OOV) problem. To address this rigidity, research has shifted toward generative models, specifically sequence-to-sequence architectures. These models treat language understanding not just as a labeling task, but as a translation-like process where the system translates a user input into a structured semantic frame [4].

The core innovation in this area involves moving beyond independent classifiers to more integrated systems. A significant contribution is the work of Zhao and Feng (2018), who proposed a generative neural network for slot filling that utilizes a Pointer-Generator mechanism [11]. This approach allows the system to decide, at each step, whether to generate a word from its internal fixed vocabulary or to copy a specific word directly from the original input.




Figure 3: Architecture of a sequence-to-sequence NLU model with a joint pointer and attention mechanism [11].

This diagram highlights the hybrid nature of the architecture proposed by Zhao and Feng. On the right, the encoder (built with GRU cells) processes the input sequence, while the attention mechanism helps the system focus on relevant parts of the text. The core breakthrough is visible in the probability formula shown in the center. The final probability of a word, P is calculated as a weighted sum of two sources: P_{gen} , which represents the likelihood of generating a word from the predefined vocabulary, and P_{ptr} , which represents the likelihood of copying a word directly from the input sequence. A learnable P_t determines the balance between these two actions. This allows the decoder

to handle unknown names or rare terms by simply copying them from the input instead of failing to find them in the vocabulary.

This dual-action strategy is particularly effective for handling the unpredictable nature of human speech. In a typical conversation, a user might mention a unique entity, such as a rare restaurant name or a specific address, which the model has never encountered before. While traditional neural networks might struggle with these unknown terms, the pointer network can simply copy them into the resulting slot without needing to categorize them first [11]. This effectively reduces the heavy reliance on delexicalization and manual feature engineering that characterizes earlier pipelines.

Furthermore, these generative architectures often incorporate attention mechanisms to maintain context over the entire sequence. As demonstrated, attention allows a model to focus on the most relevant parts of a sequence regardless of its length, helping the system stay aligned with the user objective even in complex sentences [8]. This is closely related to the joint modeling approach where a single network concurrently handles intent detection and slot filling. By sharing information across these tasks, the system can leverage the natural correlation between the overall goal and the specific details, leading to more robust performance and a deeper understanding of the semantic structure [5, 7].

By treating NLU as a translation task from human language to structured frames, this generative approach provides a balance between the deep linguistic intelligence of modern neural networks and the practical flexibility needed to handle the diverse inputs of real-world dialogue.

3.3 Overview of BERT (Ciobanu Sergiu-Tudor)

While LSTMs networks were the previous standard for NLU, they present significant limitations in training efficiency. Due to the fact that LSTMs process tokens sequentially, they require a high number of time steps to propagate information, which complicates the learning of long-range dependencies. Even bidirectional LSTMs (Bi-LSTMs) are restricted, because they capture context by processing sequences from both directions independently before concatenating the results.

The Bidirectional Encoder Representations from Transformers (BERT) [3] architecture addresses these issues by processing all words in a sequence simultaneously. This parallelization significantly increases training speed and allows the model to capture deeper contextual relationships. In the Transformer encoder, input words are passed through the architecture at once to generate embeddings that encapsulate the specific meaning of each token based on its entire surrounding context.

The BERT framework utilizes a two-stage paradigm consisting of pre-training followed by fine-tuning. During the pre-training phase, the model is trained on unlabeled corpora using two primary objectives: Masked Language Modeling (MLM) and Next Sentence Prediction (NSP).

The structural configuration of the BERT-Base model, which serves as the foundation for the methods discussed in this report, comprises 12 transformer blocks (layers), a hidden state size of 768, and 12 self-attention heads. The input representation is standardized by the inclusion of special tokens: a classification token ([CLS]) is positioned at the start of every input sequence, and a separator token ([SEP]) is utilized to denote the termination of a sequence.

3.4 JointBERT: Joint Intent Classification and Slot Filling (Ciobanu Sergiu-Tudor)

To address the limitations of independent NLU pipelines, the JointBERT model [2] proposes a unified framework for intent classification and slot filling. By leveraging the bidirectional representations of BERT, the model enhances semantic concept extraction by ensuring that the detected intent and the identified slots are mutually consistent within the same context.

JointBERT 5 is implemented by adding two task-specific, fully-connected layers directly on top of the pre-trained BERT encoder. The final hidden state of each token, denoted as $h \in \mathbb{R}^d$, is utilized for prediction as follows:

- **Intent Classification:** The model utilizes the hidden state of the first special token, $h_{[CLS]}$, as a representative embedding for the entire sequence. This vector is fed into a softmax layer to determine the intent label y^i .
- **Slot Filling:** The hidden states of the remaining tokens, $h_1, \dots, h_n$, are mapped to a softmax over the slot labels using the IOB (Inside-Outside-Beginning) tagging format. To prevent tag inconsistencies, the model only performs classification on the first subword of any token split by the BERT tokenizer.

The fundamental principle of JointBERT is the minimization of the cross-entropy loss. This allows the model to learn the interdependencies between the overall intent and the specific slot values. Following this methodology, the joint objective seeks to maximize the conditional probability of the intent y^i and the sequence of slot labels y^s given the input x :

$$P(y^i, y^s | x) = P(y^i | x) \prod_{n=1}^N P(y^{s_n} | x) \quad (1)$$




Figure 4: High-level view of the JointBERT Architecture [2].

3.5 TOD-BERT: Pre-trained NLU for Task-Oriented Dialogue (Ciobanu Sergiu-Tudor)

While BERT and its variants have achieved state-of-the-art results in general NLP benchmarks, they often demonstrate minimal improvements in conversational applications. To bridge this gap, TOD-BERT (Task-Oriented Dialogue BERT) was developed by further pre-training the BERT-Base architecture on a massive, aggregated corpus of task-specific dialogue data. This approach allows the model to better capture the underlying interaction behavior and speaker-specific information inherent in human-to-human conversations.

To tailor the model for dialogue, two special tokens, [USR] and [SYS], were added to the byte-pair embeddings and prefixed to each user utterance and system response, respectively. The training corpus consists of an aggregation of nine English, human-human, multi-turn dialogue datasets, totaling approximately 100,000 dialogues. This dataset encompasses over 1.388 million utterances across 60 distinct domains, providing a diverse foundation for learning cross-domain conversational patterns.

TOD-BERT utilizes a dual-objective learning strategy to capture both local linguistic nuances and global dialogue coherence through two specific loss functions:

1. **Masked Language Modeling (MLM):** Similar to the original BERT, tokens are randomly masked to train the model on predicting missing words based on bidirectional context. The loss is defined as:

$$L_{mlm} = - \sum_{m=1}^M \log P(x_m) \quad (2)$$

where M is the total number of masked tokens and $P(x_m)$ is the predicted probability of the token.

2. **Response Contrastive Loss (RCL):** Unlike general Next Sentence Prediction,

RCL uses a dual-encoder style to capture response similarity and sequential order. A dialogue is split into two segments (X and Y). The model employs in-batch negative training, treating the task as a ranking problem where the correct response pair (X_i, Y_i) must be distinguished from other responses in the batch. The RCL loss is given by:

$$L_{rcl} = - \sum_{i=1}^b \log M_{i,i}, \quad M = \text{Softmax}(CR^T) \in \mathbb{R}^{b \times b} \quad (3)$$

where C and R represent the encoded context and response representations, and b denotes the batch size.




Figure 5: The TOD-BERT Architecture [10].

4 Evaluating NLU Performance (Moga Patricia)

The development of reliable dialogue systems relies heavily on a strong evaluation strategy. By shifting from simple measurements toward sophisticated comparative studies, researchers can verify that a model remains stable and useful when applied to practical, everyday situations.

4.1 Quantitative Metrics

The performance of a model is typically measured through the **F-score**, which is the harmonic mean of precision and recall. This metric is essential for dialogue act prediction and slot filling tasks, as it balances the accuracy of the identified information with the completeness of the results. In addition to the F-score, the reduction in cross-entropy is used as a standard measurement for session modeling, while success and reward levels are monitored during user-in-the-loop tasks to evaluate practical utility [7].

$$F_1 = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} \quad (4)$$

Precision ensures the reliability of the extracted concepts, while recall ensures that the system captures every necessary detail from the input. These metrics provide an objective view of the model’s technical precision before it is integrated into a larger application.

4.2 Comparative Benchmarks and Simulation

In terms of performance, studies often show that models using conversation history and context do a much better job. This is especially important for longer interactions where the system needs to remember previous parts of the chat to stay relevant. By moving away from basic, fixed datasets and focusing on the actual flow of the conversation, these systems become much more accurate in extracting the right information.

To make the training process more effective, we often use sequence-based simulators to act like human users. These tools are a big improvement because they follow real human patterns much more closely than older methods, such as n-gram or agenda-based simulators. Additionally, modern multi-dimensional approaches show that they can reach the same high quality as traditional ones, but they have the advantage of being much faster to train. This makes them very useful because they require less data when they need to be adapted to a new topic or domain.

4.3 Insights and Robustness

Practical observations have led to several important insights regarding the stability of NLU models:

- **Cross-Domain Transfer:** Systems that are designed to be independent of a specific topic are very good at applying what they have learned to new areas. This flexibility makes them much more reliable when the conversation shifts to a different subject.
- **Relational Performance:** Models that focus on how different pieces of information connect work much better when the user’s input is complex. In these situations, simpler models usually lose their accuracy quite quickly, while these specialized models stay precise.
- **Simulation Fidelity:** When we give the system more detailed options for its actions, the simulated users start to behave much more like real people. This makes the whole testing process feel more natural and ensures that the evaluation results are closer to what happens in real life.

These findings demonstrate that moving toward context-aware and flexible architectures is necessary for maintaining stability as dialogue systems encounter more complex and diverse inputs.

5 Conclusion and Future Directions

The evolution of Natural Language Understanding has followed a natural and logical progression, moving from rigid, rule-based systems toward sophisticated generative and transformer-based architectures. Throughout this report, we have explored how modular pipelines provide a stable framework for processing language by breaking down interactions into intent and slot components.

Although modular architectures serve as the established foundation for traditional NLU, the industry is increasingly shifting toward end-to-end processes that utilize Large Language Models. These modern frameworks operate as a single network for everything, replacing traditional sequential architectures with an integrated system capable of managing context and generating responses at the same time.

This transition toward end-to-end systems offers significant benefits, especially the joint optimization of the entire conversation through backpropagation. This helps eliminate error accumulation. Furthermore, these models are very flexible and better at handling the unpredictable nature of real-world dialogue. However, they also have challenges. Traditional modular systems still provide better explainability and are easier to debug, whereas end-to-end networks can be less predictable and require massive amounts of high-quality data to work correctly.

Overall, the shift from simple classification to modern Large Language Models shows a clear focus on making digital interactions feel more natural. Looking ahead, the main challenge will be finding a balance between the broad intelligence of integrated systems and the high precision required for specific, task-oriented applications. Even as technology moves toward end-to-end processing, the core principles covered in this report remain essential for understanding how these complex models truly perceive and interact with the world.

References

- [1] Zefan Cai, Xin Zheng, Tianyu Liu, Haoran Meng, Jiaqi Han, Gang Yuan, Binghuai Lin, Baobao Chang, and Yunbo Cao. DialogVCS: Robust natural language understanding in dialogue system upgrade. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, Stroudsburg, PA, USA, 2024. Association for Computational Linguistics.
- [2] Qian Chen, Zhu Zhuo, and Wen Wang. Bert for joint intent classification and slot filling, 2019.
- [3] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding, 2018.
- [4] Jianfeng Gao, Michel Galley, and Lihong Li. Neural approaches to conversational AI. September 2018.
- [5] Dilek Hakkani-Tür, Gokhan Tur, Asli Celikyilmaz, Yun-Nung Chen, Jianfeng Gao, Li Deng, and Ye-Yi Wang. Multi-domain joint semantic frame parsing using bi-directional RNN-LSTM. In *Interspeech 2016*, ISCA, September 2016. ISCA.
- [6] Shengpeng Ji, Yifu Chen, Minghui Fang, Jialong Zuo, Jingyu Lu, Hanting Wang, Ziyue Jiang, Long Zhou, Shujie Liu, Xize Cheng, Xiaoda Yang, Zehan Wang, Qian Yang, Jian Li, Yidi Jiang, Jingzhen He, Yunfei Chu, Jin Xu, and Zhou Zhao. Wavchat: A survey of spoken dialogue models, 2024.
- [7] Bing Liu and Ian Lane. Attention-based recurrent neural network models for joint intent detection and slot filling. September 2016.
- [8] Colin Raffel and Daniel P W Ellis. Feed-forward networks with attention can solve some long-term memory problems. December 2015.
- [9] Owen Rambow, Srinivas Bangalore, and Marilyn Walker. Natural language generation in dialog systems. In *Proceedings of the First International Conference on Human Language Technology Research, HLT '01*, page 1–4, USA, 2001. Association for Computational Linguistics.
- [10] Chien-Sheng Wu, Steven C.H. Hoi, Richard Socher, and Caiming Xiong. Tod-bert: Pre-trained natural language understanding for task-oriented dialogue. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, page 917–929. Association for Computational Linguistics, 2020.

- [11] Lin Zhao and Zhe Feng. Improving slot filling in spoken language understanding with joint pointer and attention. In Iryna Gurevych and Yusuke Miyao, editors, *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 426–431, Melbourne, Australia, July 2018. Association for Computational Linguistics.

Acknowledgement: This work is the result of my own activity, and I confirm I have neither given, nor received unauthorized assistance for this work. I declare that I used generative AI or automated tools in the creation of content or drafting of this document and that the sequences underlined throughout this document represent the specific sections where these tools were used for assistance.